Deep-Learning-based Image Denoising in Ophthalmology
GR kick-off

Lukas Krenz
Advisor: Nicola Rieke
Director: Prof. Dr. Nassir Navab

January 12, 2018

TUM, Chair for Computer Aided Medical Procedures & Augmented Reality

Example Use-Case: Digital Window

Digital Window

What are we trying to do?

Idea

- Improve quality of retinal images
- Should work in an intraoperative setting

Use cases

- Pre-processing for computer vision algorithms (e.g. vessel segmentation)
- Provide improved quality for digital zoom.

Constraints

- Real-time
- Different levels of zoom
- Preserve structure of vessels etc.

General Super-Resolution

Many different super resolution approaches.

E.g. based on image statistics, self-simiarity of images, learned patch-databases or **convolutional neural networks**.

Only CNNs fulfil **both** quality and speed constraints

Opthamology

Use CNN-Upscaling[1]

Solid results, not optimized for intraoperative use-case

---

[1] D. Mahapatra, B. Bozorgtabar, S. Hewavitharanage and R. Garnavi (2017). 'Image Super Resolution Using Generative Adversarial Networks and Local Saliency Maps for Retinal Image Analysis'. In: *International Conference on Medical Image Computing and Computer-Assisted Intervention*.

Some common CNN-Topologies

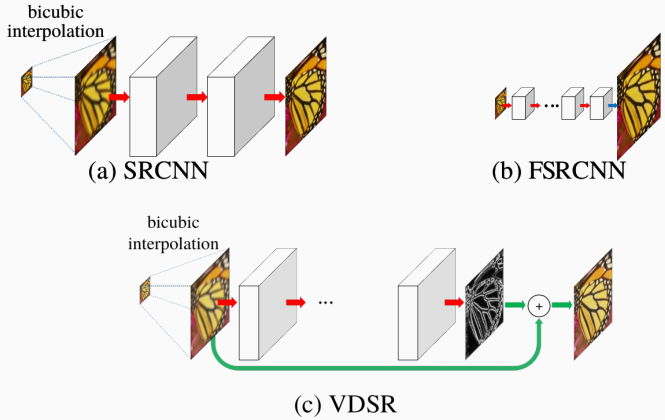

Common Topologies[2]

[2] W.-S. Lai, J.-B. Huang, N. Ahuja and M.-H. Yang (2017). 'Deep Laplacian Pyramid Networks for Fast and Accurate Super-Resolution'. In: *IEEE Conference on Computer Vision and Pattern Recognition*.

Laplacian Network[3]

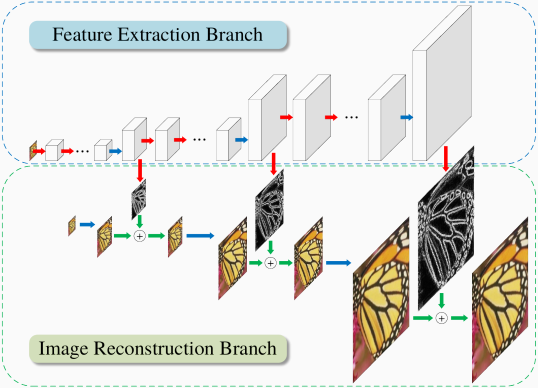

Laplacian pyramid network

1. Progressive Reconstruction
2. Residual Learning
3. Low-resolution filters
4. Non-SR-case: Add residual to HR-images

[3] Lai, Huang, Ahuja and Yang, see n. 2.

Robust Loss

(Relaxed) L_1 loss instead of L_2 loss

$$L(\hat{\boldsymbol{y}}, \boldsymbol{y}; \boldsymbol{\theta}) = \sum_n p\left(\hat{\boldsymbol{y}} - \boldsymbol{y}\right)$$

$$p(\boldsymbol{x}) = \sqrt{\langle x, x \rangle + \epsilon^2},$$

Robust, avoids **blurry** images[4]

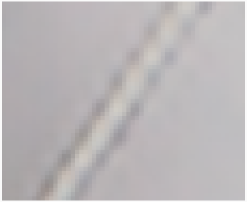

L_2 loss

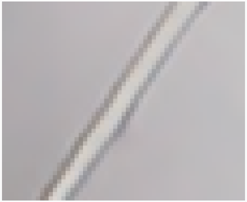

L_1 loss

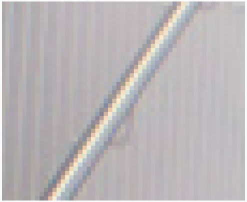

Ground truth

[4] Lai, Huang, Ahuja and Yang, see n. 2.

Perceptual Loss—VGG-based

Transform images before calculating loss, with feature map ϕ:

$$L_p(\hat{\boldsymbol{y}}, \boldsymbol{y}; \boldsymbol{\theta}) = \|\phi(\hat{\boldsymbol{y}}) - \phi(\boldsymbol{y})\|$$

Possibility 1: Neural Network based loss, **automatic** features[5]

Extract feature maps from pre-trained VGG16-Network

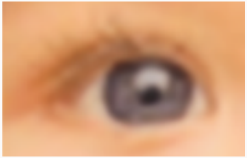

L_2 loss

Perceptual loss

Ground truth

[5] J. Johnson, A. Alahi and L. Fei-Fei (2016). 'Perceptual losses for real-time style transfer and super-resolution'. In: *European Conference on Computer Vision*.

Perceptual Loss—Saliency Maps

Possibility 2: Saliency Maps[6]

Handcrafted features

Retina image

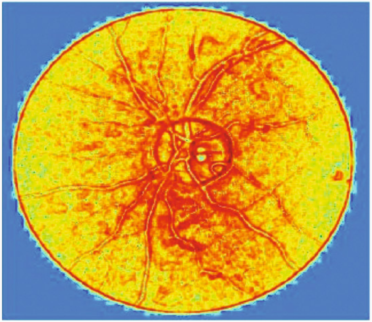

Resulting saliency map

[6] Mahapatra, Bozorgtabar, Hewavitharanage and Garnavi, see n. 1.

Adversarial Loss

Two player game between

Generator Super-resolution network

Discriminator Tries to decide, whether images are **real** high-resolution images or **generated** by our network

Leads to realistic looking images

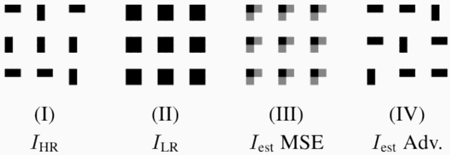

Toy example for adversarial loss[7]

[7] M. S. M. Sajjadi, B. Schölkopf and M. Hirsch (2017). 'EnhanceNet: Single Image Super-Resolution through Automated Texture Synthesis'. In: *IEEE International Conference on Computer Vision (ICCV 2017)*.

Possible Modifications for Network

Weight sharing

Replace stacked convolutions

Replace transposed convolutions (avoid checkerboard artifacts)[8]

Example of block artifacts (from[9])

[8] A. Odena, V. Dumoulin and C. Olah (2016). 'Deconvolution and Checkerboard Artifacts'. In: *Distill*.

[9] Johnson, Alahi and Fei-Fei, see n. 5.

Dataset

Use 1000 images from Eyepacs dataset[10] (ca. 300k HR) for training

Crop black borders, rescale to 1024×1024

Augment, then select random crops of size 128×128, create low-resolution images by downscaling and blurring

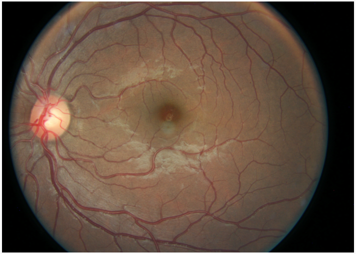

Eyepacs example

[10] `https://www.kaggle.com/c/diabetic-retinopathy-detection`

Evaluation

Normally: Compare pixel-wise reconstruction error

Useless in our case.

Perceptual & adversarial loss improve perceived quality but increase error!

Rather compare effectiveness as a pre-processing tool.

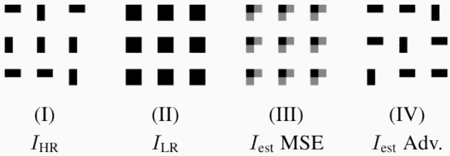

Adversarial loss higher quality with larger error[11]

[11] Sajjadi, Schölkopf and Hirsch, see n. 7.

Timeline

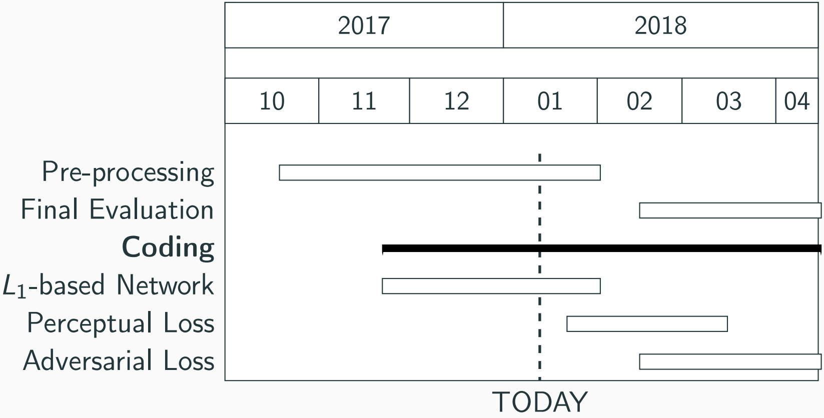
2017
2018
10
11
12
01
02
03
04
Pre-processing
Final Evaluation
Coding
L_1-based Network
Perceptual Loss
Adversarial Loss
TODAY

Summary

- Real-time super-resolution is realistic
- Selection of loss function is important
- Need for careful evaluation